

# project

September 22, 2025

## 1 Harmonic Tidal Analysis - Seattle NOAA Station

### 1.1 The project!

This short project analyses 4 months of real water level from a Seattle NOAA station to fit a harmonic tidal model and investigate outliers of the model. We use harmonic analysis to decompose the tidal data into its ‘constituent astronomical components’  $M_1, S_1, K_1$  (more on this later) and examine how well this physics centered approach fits real data!

```
[10]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

### 1.2 Loading data!

We have 4 CSVs from NOAA’s IOOS system covering May-August 2025 Seattle with 60minute interval water level measurements. Lets load to a pandas dataframe!

```
[9]: # May 2025 -> August 2025!
files = [
    "data/IOOS_Raw_Water_Level_202505.csv",
    "data/IOOS_Raw_Water_Level_202506.csv",
    "data/IOOS_Raw_Water_Level_202507.csv",
    "data/IOOS_Raw_Water_Level_202508.csv"
]

dfs = []
for file in files:
    df_month = pd.read_csv(file, parse_dates=["time"], skiprows=[1]) # row 1 is ↵
    ↵units
    dfs.append(df_month)

df = pd.concat(dfs, ignore_index = True)
df
```

```
[9]:
```

		time	WL_VALUE	latitude	longitude	STATION_ID	\
0	2025-05-01	00:00:00+00:00	3.779	47.6026	-122.3393	9447130	
1	2025-05-01	00:06:00+00:00	3.876	47.6026	-122.3393	9447130	

2	2025-05-01	00:12:00+00:00	3.974	47.6026	-122.3393	9447130
3	2025-05-01	00:18:00+00:00	4.062	47.6026	-122.3393	9447130
4	2025-05-01	00:24:00+00:00	4.153	47.6026	-122.3393	9447130
...	...	...	...	...	...	...
29514	2025-08-31	23:30:00+00:00	5.056	47.6026	-122.3393	9447130
29515	2025-08-31	23:36:00+00:00	5.049	47.6026	-122.3393	9447130
29516	2025-08-31	23:42:00+00:00	5.040	47.6026	-122.3393	9447130
29517	2025-08-31	23:48:00+00:00	5.034	47.6026	-122.3393	9447130
29518	2025-08-31	23:54:00+00:00	5.031	47.6026	-122.3393	9447130

DATUM	
0	MSL
1	MSL
2	MSL
3	MSL
4	MSL
...	...
29514	MSL
29515	MSL
29516	MSL
29517	MSL
29518	MSL

[29519 rows x 6 columns]

We can see from above there is: time of measurement, WL\_VALUE is tide height in meters, and the rest we dont need (maye for a future project...)

It seems sorted in time and without NAs but just to be sure:

```
[12]: df = df.dropna(subset=["time", "WL_VALUE"])
df["WL_VALUE"] = pd.to_numeric(df["WL_VALUE"], errors="coerce")
df = df.dropna(subset=["WL_VALUE"])
df
```

```
[12]:
```

		time	WL_VALUE	latitude	longitude	STATION_ID \
0	2025-05-01	00:00:00+00:00	3.779	47.6026	-122.3393	9447130
1	2025-05-01	00:06:00+00:00	3.876	47.6026	-122.3393	9447130
2	2025-05-01	00:12:00+00:00	3.974	47.6026	-122.3393	9447130
3	2025-05-01	00:18:00+00:00	4.062	47.6026	-122.3393	9447130
4	2025-05-01	00:24:00+00:00	4.153	47.6026	-122.3393	9447130
...	...	...	...	...	...	...
29514	2025-08-31	23:30:00+00:00	5.056	47.6026	-122.3393	9447130
29515	2025-08-31	23:36:00+00:00	5.049	47.6026	-122.3393	9447130
29516	2025-08-31	23:42:00+00:00	5.040	47.6026	-122.3393	9447130
29517	2025-08-31	23:48:00+00:00	5.034	47.6026	-122.3393	9447130
29518	2025-08-31	23:54:00+00:00	5.031	47.6026	-122.3393	9447130

	DATUM
0	MSL
1	MSL
2	MSL
3	MSL
4	MSL
...	...
29514	MSL
29515	MSL
29516	MSL
29517	MSL
29518	MSL

[29519 rows x 6 columns]

### 1.2.1 Now lets convert to time series pandas

```
[17]: s_6min = df.set_index("time")["WL_VALUE"].astype(float).sort_index()
      s_hour = s_6min.resample("1h").mean()

      print("Coverage:", s_6min.index.min(), "→", s_6min.index.max())
      print("Data points:", len(s_6min), "6-min intervals |", len(s_hour), "hourly_
        ↳points")

      print("\nData quality checks:")

      # 6 minute interval check
      diff = s_6min.index.to_series().diff().value_counts().head()
      print("Top time intervals:")
      print(diff)
      print("(Should be mostly 00:06:00 for 6-minute data)")

      # Make sure all data is from same station and datum
      print(f"\nStations: {df['STATION_ID'].dropna().unique()}")
      print(f"Datum: {df['DATUM'].dropna().unique()}")
      print("(Should be single station and datum for consistent analysis)")
```

Coverage: 2025-05-01 00:00:00+00:00 → 2025-08-31 23:54:00+00:00

Data points: 29519 6-min intervals | 2952 hourly points

Data quality checks:

Top time intervals:

time

0 days 00:06:00      29517

0 days 00:12:00          1

Name: count, dtype: int64

(Should be mostly 00:06:00 for 6-minute data)

Stations: [9447130]

Datum: ['MSL']

(Should be single station and datum for consistent analysis)


1.3 okay looks all good (with one 12 minute gap) Let's plot the tide vs date:

```
[20]: print(f"Water level range: {s_6min.min():.3f} to {s_6min.max():.3f} meters")
      print(f"Mean: {s_6min.mean():.3f}, Std: {s_6min.std():.3f}")
```

```
plt.figure(figsize=(15, 6))
# using pandas time series .plot
s_6min.plot(title="Full 6-minute Water Level Time Series - Seattle NOAA_
↳Station", fontsize= 16)
plt.ylabel("Water Level (meters)", fontsize = 12)
plt.xlabel("Time", fontsize=12)
plt.grid(True, alpha=0.3)
plt.show()
```

Water level range: 1.215 to 6.247 meters

Mean: 4.459, Std: 1.124



1.4 cool! Can see the sinusoidal patterns, Let's group by time of Day:

```
[23]: df_analysis = s_6min.reset_index()
      df_analysis['hour'] = df_analysis['time'].dt.hour
      df_analysis['minute'] = df_analysis['time'].dt.minute
      df_analysis['time_of_day'] = df_analysis['hour'] + df_analysis['minute']/60

      time_stats = df_analysis.groupby('time_of_day')['WL_VALUE'].agg([
```

```

    'count', 'median', 'mean', 'std',
    lambda x: x.quantile(0.05), # 5th percentile
    lambda x: x.quantile(0.95) # 95th percentile
])).round(3)

time_stats.columns = ['count', 'median', 'mean', 'std', 'p5', 'p95']

print("Sample of time-of-day statistics:")
print(time_stats.head(10))
print(f"\nWe have {len(time_stats)} unique times of day (should be 240 for 6-min intervals)")

plt.figure(figsize=(12, 6))

plt.scatter(df_analysis['time_of_day'], df_analysis['WL_VALUE'],
            c='lightgray', alpha=0.8, s=1.9, label='All measurements')

plt.plot(time_stats.index, time_stats['median'], 'ro-',
         linewidth=2, markersize=3, label='Median by time of day')


plt.title('Daily Tidal Pattern - All Data Points + Median Line')
plt.xlabel('Hour of Day')
plt.ylabel('Water Level (meters)')
plt.xlim(0, 24)
plt.xticks(range(0, 25, 3)) # Every 3 hours
plt.grid(True, alpha=0.3)
plt.legend()
plt.tight_layout()
plt.show()

```

Sample of time-of-day statistics:

	count	median	mean	std	p5	p95
time_of_day						
0.0	123	4.776	4.561	0.988	2.745	5.796
0.1	123	4.769	4.594	0.980	2.760	5.788
0.2	123	4.800	4.626	0.972	2.769	5.788
0.3	123	4.832	4.658	0.964	2.780	5.812
0.4	123	4.846	4.689	0.956	2.827	5.830
0.5	123	4.846	4.720	0.948	2.887	5.842
0.6	123	4.933	4.751	0.939	2.932	5.855
0.7	123	4.951	4.780	0.931	2.923	5.870
0.8	123	4.970	4.809	0.923	2.918	5.882
0.9	123	5.023	4.837	0.914	2.973	5.885

We have 240 unique times of day (should be 240 for 6-min intervals)




1.5 Very cool! 1 More plot before fitting... Lets do a zoom in on the first two weeks:

```
[25]: # Let's first plot a shorter time window to see the tidal patterns clearly
plt.figure(figsize=(15, 6))

# Plot just 2 weeks to see the patterns
two_weeks = s_hour.head(24*14) # First 14 days
plt.plot(two_weeks.index, two_weeks.values, 'b-', linewidth=1)

plt.title('Water Level - First 2 Weeks (to see tidal patterns)')
plt.xlabel('Time')
plt.ylabel('Water Level (meters)')
plt.grid(True, alpha=0.3)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



## 1.6 Sure Looks like a sum of sinuisoidal curves!

Now for the physics.

### 1.6.1 Harmonic Tidal Thoery

Tides can be modeled from a sum of sine waves where each wave represents a differnet astronomical ‘forcing’ - the moon and sun pulling on oceans from different angles and distances.

The water level can be modeled using these sinusoidal components like:

$$\text{Water Level}(t) = \text{Mean} + \sum_i A_i \cos\left(\frac{2\pi}{T_i}t + \phi_i\right)$$

Where  $T_i$  is the period of each contituent,  $A_i$  is the amplitude, and  $\phi_i$  is the phase.

There are many constituents, but we will the following 6 that have short periods that can be fit with our 4-month data and are known(?) to be strong in the Pacific Coast. Also Diurnal - one period per day. Semi-diurnal twice per day

**The 6 major constituents come from:** - **M2 (12.42h)**: Principal lunar semi-diurnal - Earth’s rotation relative to moon - **S2 (12.00h)**: Principal solar semi-diurnal - Earth’s rotation relative to sun

- **N2 (12.66h)**: Lunar elliptic semi-diurnal - moon’s elliptical orbit variations - **K1 (23.93h)**: Lunar diurnal - moon’s declination (north-south position) - **O1 (25.82h)**: Lunar diurnal - combined moon declination and orbit effects - **P1 (24.07h)**: Solar diurnal - sun’s declination variations

Taken from page 98 of [https://tidesandcurrents.noaa.gov/publications/Tidal\\_Analysis\\_and\\_Predictions.pdf](https://tidesandcurrents.noaa.gov/publications/Tidal_Analysis_and_Predictions.pdf)

Theese periods are fixed by celesital mechanics. Only  $A_i$  and  $\phi_i$  very by location.

So essentially we will be summing the following sinuisodal and fitting the amplitude  $A_i$  and  $\phi_i$  to the Seattle Tidal data while fixing the  $T_i$  to the NOAA values. The following plots the periods of the 6 major Tidal Constituents

```
[26]: fig, axes = plt.subplots(3, 2, figsize=(15, 10))
fig.suptitle('The 6 Major Tidal Constituents (Period compairons)', fontsize=16)

t_hours = np.linspace(0, 72, 1000) # 3 days

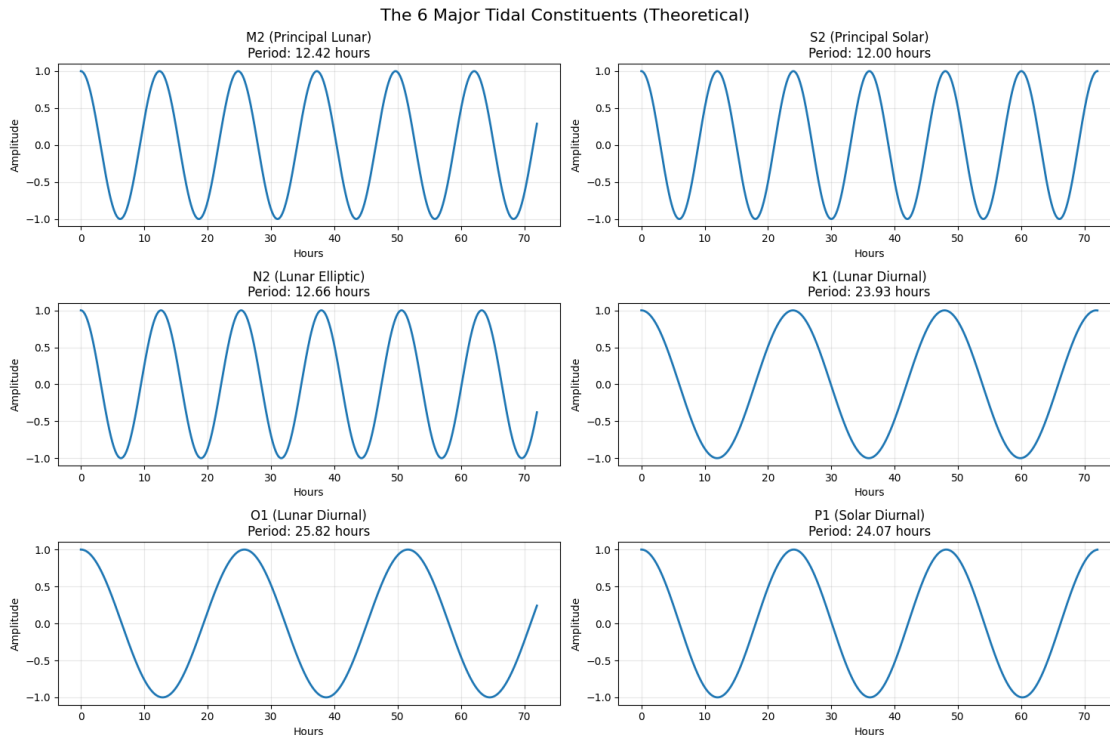
constituents = {
    'M2 (Principal Lunar)': 12.421,
    'S2 (Principal Solar)': 12.000,
    'N2 (Lunar Elliptic)': 12.658,
    'K1 (Lunar Diurnal)': 23.934,
    'O1 (Lunar Diurnal)': 25.819,
    'P1 (Solar Diurnal)': 24.066
}

for i, (name, period) in enumerate(constituents.items()):
    row, col = i // 2, i % 2
    omega = 2 * np.pi / period
    wave = np.cos(omega * t_hours)

    axes[row, col].plot(t_hours, wave, linewidth=2)
    axes[row, col].set_title(f'{name}\nPeriod: {period:.2f} hours')
    axes[row, col].set_xlabel('Hours')
    axes[row, col].set_ylabel('Amplitude')
    axes[row, col].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```





## 1.7 Okay now for the fit:

We will use `scipy`'s `curve_fit` to fit the amplitude and phase of the 6 constituents as well as the mean tidal height per hour leaving 13 parameters.

Specifically `curve_fit` uses least squares optimization or finds the 13 parameters that minimize the sum of squared differences between model predictions and the **hourly averaged** water levels. Each hour's data point is an average of ~10 six-minute measurements from that hour noise

```
[32]: from scipy.optimize import curve_fit

# 6-constituent harmonic model
def tidal_model(t_hours, mean, A_M2, phi_M2, A_S2, phi_S2, A_N2, phi_N2,
                A_K1, phi_K1, A_O1, phi_O1, A_P1, phi_P1):
    """
    Harmonic tidal model with 6 constituents
    13 parameters total: 1 mean + 6 amplitudes + 6 phases
    """
    periods = {
        'M2': 12.421, 'S2': 12.000, 'N2': 12.658,
        'K1': 23.934, 'O1': 25.819, 'P1': 24.066
    }

    # to angular frequencies
```

```

omega_M2 = 2 * np.pi / periods['M2']
omega_S2 = 2 * np.pi / periods['S2']
omega_N2 = 2 * np.pi / periods['N2']
omega_K1 = 2 * np.pi / periods['K1']
omega_O1 = 2 * np.pi / periods['O1']
omega_P1 = 2 * np.pi / periods['P1']

# mean, sum of all constituents
result = (mean +
          A_M2 * np.cos(omega_M2 * t_hours + phi_M2) +
          A_S2 * np.cos(omega_S2 * t_hours + phi_S2) +
          A_N2 * np.cos(omega_N2 * t_hours + phi_N2) +
          A_K1 * np.cos(omega_K1 * t_hours + phi_K1) +
          A_O1 * np.cos(omega_O1 * t_hours + phi_O1) +
          A_P1 * np.cos(omega_P1 * t_hours + phi_P1))

return result

# using hours
t0 = s_hour.index[0]
t_hours = (s_hour.index - t0).total_seconds() / 3600
water_levels = s_hour.values

# Initial parameter guess: [mean, A1, phi1, A2, phi2, ...]
p0 = [s_hour.mean(), # mean
      0.5, 0, # M2 amplitude, phase
      0.3, 0, # S2 amplitude, phase
      0.2, 0, # N2 amplitude, phase
      0.4, 0, # K1 amplitude, phase
      0.3, 0, # O1 amplitude, phase
      0.1, 0] # P1 amplitude, phase

print("Starting harmonic fit...")
print(f>Data points: {len(water_levels)}")
print(f>Time span: {t_hours[-1]:.1f} hours ({t_hours[-1]/24:.1f} days)")

# fit the model
popt, pcov = curve_fit(tidal_model, t_hours, water_levels, p0=p0)

print("Fit completed!")
print(f>Mean water level: {popt[0]:.3f} meters")

```

```

Starting harmonic fit...
Data points: 2952
Time span: 2951.0 hours (123.0 days)
Fit completed!
Mean water level: 4.456 meters

```

```
[ ]:
```

Now that the fit is complete I will make a figure of 3 plots.

- **Top plot** is the observed water level (hourly average) vs the predicted harmonic fit water level for the first 30 days to see how well the model captures tidal cycles.
- **Middle Plot** shows the residual between the hourly averaged tidal heights and model predictions over the full 4 months. Important to look for patterns in the residuals demonstrating what the tidal
- **Bottom plot:** Bar chart of the fitted amplitudes for each of the 6 tidal constituents, showing relative importance

```
[33]: fitted_values = tidal_model(t_hours, *popt)
residuals = water_levels - fitted_values

fig, axes = plt.subplots(3, 1, figsize=(15, 12))

# plot 1: Observed vs Fitted (first 30 days)
days_to_show = 30
end_idx = min(24 * days_to_show, len(s_hour))
time_subset = s_hour.index[:end_idx]
obs_subset = water_levels[:end_idx]
fit_subset = fitted_values[:end_idx]

axes[0].plot(time_subset, obs_subset, 'b-', alpha=0.7, linewidth=1,
             label='Observed')
axes[0].plot(time_subset, fit_subset, 'r-', linewidth=1.5, label='Harmonic Fit')
axes[0].set_title(f'Tidal Harmonic Fit - First {days_to_show} Days')
axes[0].set_ylabel('Water Level (m)')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# plot 2: Residuals (observed - fitted)
axes[1].plot(s_hour.index, residuals, 'g-', alpha=0.6, linewidth=0.5)
axes[1].axhline(y=0, color='black', linestyle='--', alpha=0.5)
axes[1].set_title('Residuals (Observed - Fitted)')
axes[1].set_ylabel('Residual (m)')
axes[1].grid(True, alpha=0.3)

# plot 3: Constituent amplitudes
constituent_names = ['M2', 'S2', 'N2', 'K1', 'O1', 'P1']
amplitudes = [popt[1], popt[3], popt[5], popt[7], popt[9], popt[11]]
axes[2].bar(constituent_names, amplitudes, color=['skyblue']*3 + ['orange']*3)
axes[2].set_title('Fitted Tidal Constituent Amplitudes')
axes[2].set_ylabel('Amplitude (m)')
axes[2].grid(True, alpha=0.3)

plt.tight_layout()
```

```
plt.show()
```

```
rmse = np.sqrt(np.mean(residuals**2))  
print(f"Root Mean Square Error: {rmse:.3f} meters")  
print(f"R2 (explained variance): {1 - np.var(residuals)/np.var(water_levels):.  
↪3f}")
```



Root Mean Square Error: 0.152 meters

R<sup>2</sup> (explained variance): 0.982

## 1.8 Results and Analysis

From the three plots we can see:

**Top plot (Fit Quality):** The harmonic model captures the main tidal oscillations reasonably well, but there appears to be some systematic mismatch - sinusoidal drift over the 30 day period.

**Middle plot (Residuals):** The residuals clearly show they are NOT random noise around zero. Instead, there's a distinct sinusoidal pattern with what appears to be a ~30-day period. This systematic pattern suggests our 6-constituent model is missing longer-period tidal components.

**Bottom plot (Constituent Strengths):** M2 (principal lunar) dominates at ~1.0m amplitude, followed by K1 (lunar diurnal) at 0.9m magnitude. This confirms Seattle has mixed semi-diurnal tides - both twice-daily and once-daily components are significant.

**The smoking gun:** That organized residual pattern suggests we're missing monthly and fortnightly tidal constituents like: - **Mf (14.77 days):** Fortnightly tide from lunar orbit variations - **Mm (27.55 days):** Monthly tide from lunar distance changes

With only 4 months of data, we have limited ability to resolve these longer periods, but the residual pattern suggests they're present and measurable. Future work should focus on gathering data from a larger time window to feel these additional constituents, and possibly compare other sites geographically close to Seattle

**Model Performance:** The fit statistics show reasonable results with RMSE of 0.152 meters and  $R^2$  of 0.982, indicating our 6-constituent model explains 98.2% of the tidal variation despite missing the longer-period components.

[ ]:

[ ]: